

Empirical Return on Investment (ROI) and Systems Governance of Enterprise Generative AI Adoption

Aryaman Dev

Institute for Econometric AI Policy & Enterprise Risk Governance

researcher@institute.org



Abstract—Enterprise adoption of Generative Artificial Intelligence (GenAI) and autonomous multi-agent systems has accelerated rapidly across enterprise software engineering, customer experience, and data operations [1], [2]. However, organizations face significant friction in measuring causal return on investment (ROI), managing token compute budgets, and governing operational risk [3], [4]. In this comprehensive review, we synthesize quantitative findings from enterprise deployments [5], [6], establish a multi-tier econometric measurement framework [7], and evaluate total cost of ownership (TCO) scaling dynamics across cloud infrastructure environments [8], [9].

Enterprise adoption of Generative Artificial Intelligence (GenAI) and autonomous multi-agent systems has accelerated rapidly across enterprise software engineering, customer experience, and data operations. However, organizations face significant friction in measuring causal return on investment (ROI), managing token compute budgets, and governing operational risk. In this comprehensive review, we synthesize quantitative findings from enterprise deployments, establish a multi-tier econometric measurement framework, and evaluate total cost of ownership (TCO) scaling dynamics across cloud infrastructure environments.

Keywords— Generative AI, Empirical Evaluation, AI Systems, Enterprise Operations, Systematic Review.

Enterprise adoption of Generative Artificial Intelligence (GenAI) has transitioned rapidly from experimental proof-of-concept pilots to mission-critical operational deployments. As organizations deploy Large Language Models (LLMs) and autonomous multi-agent systems across software engineering, customer experience, and business intelligence, quantifying return on investment (ROI) has emerged as a central strategic requirement [4].

However, measuring GenAI ROI presents complex methodological challenges. Traditional marketing mix modeling (MMM) and multi-touch attribution (MTA) frameworks operate in silos, failing to capture the non-linear interaction between practitioner domain mastery, tool integration depth, and autonomous agent capabilities [7], [10], [11]. Furthermore, high compute costs, GPU memory constraints, model drift, and data security risks threaten to erode projected financial gains unless mitigated by mature MLOps governance [4], [12]–[14].

This paper delivers a principal-level literature review and empirical framework addressing:

- 1) *Enterprise Risk Governance: Operationalizing risk management boundaries to control shadow AI API provisioning, model drift, and proprietary data leakage.*

1 QUANTITATIVE ANALYSIS & EMPIRICAL EVIDENCE

The assessment of Return on Investment (ROI) for Generative AI (GenAI) initiatives in enterprise settings remains a nascent yet critical area of research. While the transformative potential of GenAI is widely acknowledged, concrete, universally comparable quantitative evidence across diverse industries is still emerging. This section presents a meta-analysis of the available quantitative results and proposed frameworks from surveyed literature, identifying key empirical findings and highlighting the methodologies being developed to substantiate GenAI's business value.

1.1 Empirical Findings: Early Indicators of Value

Quantitative evidence directly attributing specific monetary value to GenAI implementations is currently limited but highly indicative of significant potential. One notable instance of reported business value comes from a practical application demonstrated in a GitHub repository:

This single significant data point, though isolated, serves as an empirical anchor, illustrating that substantial financial benefits are attainable. It likely encompasses a combination of revenue uplift, cost savings, and efficiency gains across various business processes.

1.2 Frameworks for Causal ROI Attribution

Beyond isolated figures, several papers focus on developing robust methodologies and frameworks for attributing and optimizing GenAI's impact, particularly in complex domains like life sciences marketing [4], [15], [16].

Kumar (2026) introduces a multi-layered Causal ROI Framework specifically tailored for the life sciences industry, addressing the limitations of traditional ROI attribution models such as Marketing Mix Modelling (MMM) and

Multi-Touch Attribution (MTA) [4]. These conventional approaches often operate in silos, leading to fragmented or contradictory signals that hinder unified decision-making. The proposed framework aims to:

- 1) *Attribute commercial spend across the Healthcare Professional (HCP) engagement journey.*
- 2) *Optimize marketing budget allocation and HCP targeting.*
- 3) *Activate GenAI-driven personalization strategies across digital and offline channels.*

This framework moves beyond simple correlation, striving for a causal understanding of GenAI's impact [4], [17], [18]. The general principle of ROI is typically defined as:

$$\text{ROI} = \frac{\text{Net Profit attributable to GenAI}}{\text{Cost of GenAI Investment}} \times 100\% \quad (1)$$

$$\text{ROI} = \frac{(\Delta R + \Delta C) - I}{I} \times 100\% \quad (2)$$

Similarly, **Modi (2026)** focuses on "Measuring Business ROI of Generative AI Adoption on Azure Cloud Platforms," indicating the development of platform-specific methodologies for ROI assessment [7]. This highlights a trend towards integrating ROI measurement capabilities directly within cloud ecosystems, potentially leveraging platform-native analytics and financial reporting tools to quantify GenAI's impact on resource utilization, operational efficiency, and revenue generation tied to cloud services.

Thukral et al. (2023), while not providing explicit quantitative figures, strongly emphasizes the potential for "concrete ROI" and "reduced risk" through customer journey optimization using LLMs [4]. Their work outlines best practices for deploying GenAI in marketing and customer experience, implicitly linking successful implementation to quantifiable outcomes such as:

- **Increased Conversion Rates:** Through personalized content and proactive engagement.
- **Reduced Customer Service Costs:** Via intelligent chatbots and self-service solutions.
- **Improved Customer Lifetime Value (CLTV):** By fostering stronger relationships and tailored experiences.
- **Accelerated Time-to-Market:** For new marketing campaigns and content creation.

These qualitative discussions underscore the various levers through which GenAI is expected to drive financial returns, guiding organizations in identifying high-value use cases.

1.3 Comparison of Reported ROI & Measurement Approaches

Given the nascent stage of robust, publicly available quantitative data, a direct statistical meta-analysis with pooled effect sizes is not yet feasible. Instead, we summarize the

identified contributions regarding GenAI ROI in Table 1, categorizing them by the type of evidence or framework presented.

Table 1: Summary of GenAI ROI Contributions and Measurement Approaches

1.4 Statistical Summaries and Identified Gaps

The meta-analysis reveals a significant gap between the theoretical recognition of GenAI's ROI potential and the widespread availability of detailed, transparent empirical studies. The existing literature largely focuses on:

- **Framework Development:** Establishing methods for *how* to measure ROI, particularly addressing causality and attribution challenges.
- **Qualitative Identification:** Pinpointing areas where GenAI is expected to deliver value (e.g., customer experience, operational efficiency).
- **Initial Practical Demonstrations:** Showcasing specific instances of value realization, often without comprehensive methodological detail for external replication or generalization.

1.5 Limitations, Applicability Boundaries, and Threats to Validity:

- 1) *Heterogeneity of Use Cases: GenAI is applied across a vast spectrum of enterprise functions (marketing, customer service, R&D, operations), making direct comparisons of ROI challenging due to varying baselines, investment scales, and metrics.*
- 2) *Attribution Complexity: Isolating the precise financial impact of GenAI from other concurrent digital transformation initiatives, changes in market conditions, or other technological interventions is inherently difficult. Causal frameworks like Kumar's are crucial here but require sophisticated data and analytical capabilities.*
- 3) *Proprietary Nature of Data: Many enterprises consider their GenAI deployment details and ROI figures as competitive intelligence, limiting public disclosure and hindering academic meta-analysis.*
- 4) *Early Stage of Adoption: The rapid evolution of GenAI technology means that many implementations are still in pilot or early deployment phases, where long-term ROI may not yet be fully realized or measured.*

1.6 Future Directions for Empirical Research

To build a more robust quantitative understanding of enterprise GenAI ROI, future research should focus on:

- **Standardized Reporting:** Encouraging enterprises to adopt standardized metrics and reporting frameworks for GenAI ROI, enabling better comparability.
- **Longitudinal Studies:** Tracking GenAI implementations over extended periods to capture long-term value, including indirect benefits and compounded effects.
- **Cross-Industry Benchmarking:** Developing industry-specific benchmarks for GenAI ROI

to provide context for individual enterprise performance.

- **Case Studies with Detailed Methodology:** Publishing comprehensive case studies that detail the investment, methodology for ROI calculation, and specific outcomes, allowing for validation and generalization.
- **Impact on Human Capital:** Quantifying the ROI derived from augmenting human capabilities, such as increased employee productivity, faster skill acquisition, and enhanced decision-making.

[7]

The successful integration and sustained value generation from Generative AI (GenAI) in an enterprise setting extend far beyond model development, critically depending on robust systems and infrastructure. Neglecting these practical concerns can erode potential return on investment (ROI), transform promising initiatives into costly liabilities, and impede scalability. This section analyzes key systems-level considerations, including compute costs, scalability requirements, deployment bottlenecks, governance frameworks, and broader organizational implementation challenges crucial for realizing enterprise GenAI ROI.

2 COMPUTE COSTS AND RESOURCE MANAGEMENT

The computational demands of GenAI models, particularly Large Language Models (LLMs), represent a significant component of the total cost of ownership (TCO). These costs are multifaceted, encompassing both model training and inference.

Training Costs: Developing or extensively fine-tuning proprietary GenAI models often requires substantial investments in Graphics Processing Units (GPUs) or specialized AI accelerators. While enterprises may opt for pre-trained models and fine-tune them, even this process can be resource-intensive, particularly for large datasets and complex architectures. Cloud platforms, such as Azure, offer scalable compute resources for GenAI adoption, allowing organizations to manage fluctuating demands and potentially measure business ROI through their offerings [19]. However, the sheer scale of modern models implies that even fractional usage can accumulate substantial cloud billing.

Inference Costs: Once deployed, the ongoing inference—the process of using the model to generate outputs—becomes the primary operational cost driver. This cost is directly proportional to the volume of requests and the complexity of the model. For enterprises integrating GenAI into high-volume customer interaction points, such as customer journey optimization [7], even small per-query costs can quickly escalate. Key factors influencing inference costs include:

- **Model Size and Architecture:** Larger models require more memory and computational cycles.
- **Query Latency Requirements:** Real-time applications demand dedicated, high-performance infrastructure.
- **Throughput:** The number of simultaneous requests the system must handle.

- **Cloud vs. On-Premise Deployment:** Cloud solutions offer flexibility and elasticity but often come with higher per-unit costs, whereas on-premise solutions demand significant upfront capital expenditure and maintenance.

To mitigate compute costs, enterprises must strategically evaluate model selection, deployment architecture, and optimization techniques. This includes leveraging smaller, more specialized models where appropriate, employing techniques such as quantization, pruning, and knowledge distillation to reduce model size and inference time, and adopting efficient serving frameworks. Hybrid cloud strategies, where sensitive or high-volume inference occurs on optimized on-premise hardware and burstable workloads leverage cloud resources, can also be considered.

$$C_{op} = N_{req} \times (C_{inference} + C_{transfer}) + C_{infrastructure} + C_{storage} \quad (3)$$

3 SCALABILITY AND PERFORMANCE ENGINEERING

Enterprise GenAI solutions must be designed for scalability from inception to accommodate increasing user loads, data volumes, and expanding use cases. A solution that performs well in a pilot phase with limited users may collapse under enterprise-wide adoption.

Horizontal and Vertical Scaling:

- **Horizontal Scaling:** Involves adding more machines (e.g., GPU instances) to distribute the workload. This is often preferred for GenAI inference, allowing systems to handle many concurrent requests. Load balancers and container orchestration platforms (like Kubernetes) are essential for managing horizontal scaling.
- **Vertical Scaling:** Involves upgrading the resources of a single machine (e.g., adding more powerful GPUs, increasing RAM). This has limits but can be effective for handling very large models that require significant memory on a single device.

Data Pipeline Scalability: GenAI applications are intensely data-driven. Scalable data ingestion, processing, storage, and retrieval pipelines are critical, especially for RAG (Retrieval-Augmented Generation) architectures that depend on up-to-date and extensive knowledge bases. Ensuring these pipelines can handle vast amounts of unstructured and structured data efficiently is paramount.

Model Serving Infrastructure: Low-latency and high-throughput model serving are critical for user experience and business outcomes. This necessitates robust MLOps practices, including automated deployment, canary releases, rollback capabilities, and continuous monitoring of model performance and resource utilization. Caching mechanisms, edge deployments, and Content Delivery Networks (CDNs) can further optimize response times for geographically dispersed users. The selection of an appropriate deployment

strategy is a key technical consideration to avoid pitfalls and ensure concrete ROI [7].

4 DEPLOYMENT BOTTLENECKS AND MLOPS MATURITY

Moving GenAI from proof-of-concept to production often reveals significant deployment bottlenecks. These typically stem from a lack of MLOps maturity within the organization.

Integration with Legacy Systems: Many enterprises operate complex IT landscapes with deeply entrenched legacy systems. Integrating new GenAI services, often built with modern microservice architectures, into these existing environments can be challenging. Data formats, API compatibility, authentication mechanisms, and operational workflows must be carefully aligned.

Data Readiness and Quality: GenAI models are highly sensitive to the quality and relevance of their input data. Data silos, inconsistent data formats, and poor data governance can create significant hurdles during deployment. Preparing enterprise data for GenAI applications often involves extensive data engineering, cleaning, transformation, and semantic enrichment.

Model Lifecycle Management: Effective MLOps ensures that GenAI models are not static assets but continuously evolving components. This involves:

- **Version Control:** Tracking model versions, associated code, and training data.
- **Automated Testing:** Ensuring model integrity and performance before deployment.
- **Continuous Monitoring:** Tracking model drift, performance degradation, and anomalous behavior in production.
- **Retraining and Redeployment:** Establishing pipelines for updating models with fresh data to maintain relevance and accuracy.
- **Security:** Safeguarding models and data from adversarial attacks, data leakage, and unauthorized access.

Failing to address these bottlenecks can lead to “technical pitfalls” that hinder successful GenAI adoption and ROI realization [7].

5 GOVERNANCE AND ETHICAL AI CONSIDERATIONS

Governance for enterprise GenAI extends beyond technical implementation to encompass ethical, legal, and compliance dimensions. These considerations are particularly critical in sensitive domains like life sciences, where precise targeting and data handling are paramount [4].

Responsible AI Principles: Enterprises must establish clear guidelines for developing and deploying GenAI systems responsibly. This includes ensuring:

- **Fairness:** Models do not perpetuate or amplify biases present in training data.
- **Transparency and Explainability:** Understanding how models arrive at their outputs, especially in critical decision-making contexts.

- **Privacy and Security:** Protecting sensitive personal and corporate data used by and generated by GenAI models.
- **Accountability:** Defining clear lines of responsibility for model outputs and their consequences.
- **Safety:** Preventing models from generating harmful, misleading, or inappropriate content.

Data Governance: Strict data governance policies are essential for GenAI. This covers:

- **Data Lineage:** Tracking the origin and transformations of all data used for training and inference.
- **Access Control:** Limiting who can access and modify sensitive data.
- **Compliance:** Adhering to regulations such as GDPR, HIPAA, and industry-specific mandates.
- **Output Validation:** Implementing human-in-the-loop processes or automated checks to validate GenAI outputs before public release or critical use.

Model Governance: This involves establishing frameworks for model validation, risk assessment, and continuous auditing. For example, a model risk management framework might categorize GenAI applications by their potential impact and prescribe corresponding levels of scrutiny and oversight. These ethical and process concerns are identified as critical for successful deployment and risk reduction [7].

6 ORGANIZATIONAL IMPLEMENTATION CHALLENGES

Beyond technical infrastructure, human and organizational factors significantly impact GenAI adoption and ROI.

Skills Gap: A widespread shortage of skilled AI engineers, data scientists, and MLOps specialists poses a significant barrier. Enterprises must invest in upskilling existing talent or acquiring new expertise to build and manage GenAI capabilities effectively.

Change Management and User Adoption: Introducing GenAI solutions often requires changes to existing workflows and job roles. Resistance to change, lack of understanding, or mistrust in AI systems can hinder adoption. Effective change management strategies, including clear communication, training, and demonstrating tangible benefits, are crucial. Understanding “which initiatives and opportunities to begin with” and ensuring the organization adapts to new capabilities is key [7].

Cross-functional Collaboration: Successful GenAI initiatives require close collaboration between business stakeholders, IT, legal, and AI teams. Business leaders must articulate clear use cases and expected outcomes, while technical teams provide realistic assessments of capabilities and limitations.

Defining and Tracking ROI: Quantifying the ROI of GenAI can be complex. Traditional ROI attribution models may fall short, necessitating new frameworks that account for both direct cost savings and indirect benefits like increased innovation, enhanced customer experience, or accelerated time-to-market [4]. Developing robust tracking mechanisms and dashboards, such as enterprise-grade PowerBI solutions for GenAI business value, are critical

for demonstrating measurable outcomes [20]. Without clear metrics and a framework for measuring success, initiatives risk losing executive support.

Strategic Alignment: Organizations must align GenAI deployments with overarching business strategy. The initial impulse to “get in the game” must be tempered by strategic planning that identifies high-value use cases and considers the organization’s current AI maturity and resource levels [7]. A phased approach, starting with well-defined pilots and iteratively expanding, can mitigate risks and build internal expertise.

7 CONCLUSION

The realization of substantial enterprise GenAI ROI is inextricably linked to robust systems and infrastructure considerations. Addressing compute costs through judicious model selection and optimization, engineering for scalability, overcoming deployment bottlenecks with mature MLOps practices, establishing comprehensive governance frameworks, and navigating organizational challenges are not merely technical tasks but strategic imperatives. A holistic approach that integrates these concerns from initial ideation through continuous operation is essential for transforming the promise of GenAI into tangible, sustainable business value.

[7]

8 CRITICAL LIMITATIONS & REVIEWER AUDIT

The burgeoning interest in Generative AI (GenAI) within the enterprise landscape necessitates a rigorous evaluation of its Return on Investment (ROI). While early indicators suggest promising avenues for value creation, a comprehensive academic understanding requires acknowledging the critical limitations inherent in current methodologies, addressing open problems, confronting data quality challenges, and anticipating potential reviewer objections and ethical considerations. This section critically examines these facets, aiming to provide a balanced perspective on the current state and future directions of enterprise GenAI ROI assessment.

9 METHODOLOGICAL LIMITATIONS IN ROI ATTRIBUTION

Measuring the true ROI of GenAI adoption in complex enterprise environments presents significant methodological hurdles. A primary challenge lies in establishing a clear causal link between GenAI interventions and observed business outcomes.

9.0.1 Challenges in Causal Attribution

Traditional ROI attribution models, such as Marketing Mix Modeling (MMM) and Multi-Touch Attribution (MTA), often fall short when attempting to isolate the precise impact of GenAI initiatives. As highlighted by Kumar (2026), these approaches frequently operate in isolation, yielding fragmented or even contradictory signals that impede unified decision-making. The introduction of GenAI adds another layer of complexity, making it difficult to disentangle its

effects from concurrent marketing campaigns, operational improvements, or external market forces. A robust causal framework, as proposed by Kumar (2026) for the life sciences, is essential but remains nascent in broader enterprise GenAI contexts. Without proper causal inference, there is a risk of over-attributing positive outcomes to GenAI, leading to inflated ROI claims.

$$Y = \alpha + \tau X_{\text{GenAI}} + \sum_{i=1}^k \eta_i Z_i + \epsilon \quad (4)$$

9.0.2 Defining and Quantifying “Value”

The definition of “value” in the context of GenAI extends beyond simple financial metrics. While cost savings (e.g., reduced operational expenses) and revenue generation (e.g., increased sales from personalized recommendations) are tangible, many benefits are intangible, such as enhanced customer experience, accelerated innovation, improved employee productivity, and better decision-making capabilities (Thukral et al., 2023). Quantifying these “soft” benefits into monetary terms for ROI calculation is notoriously difficult and often relies on proxy metrics or subjective assessments, introducing potential biases.

9.0.3 Generalizability and Context Dependency

Many reported GenAI ROI figures stem from specific case studies or pilot programs (e.g., Modi, 2026; Kapoor, 2025). The success observed in one particular cloud environment (e.g., Azure platforms), industry (e.g., life sciences), or for a specific use case (e.g., customer journey optimization) may not be directly transferable to other enterprise contexts. Factors such as organizational maturity in AI adoption, existing technological infrastructure, data governance policies, and employee skill sets significantly influence deployment success and, consequently, ROI. Generalizing specific success stories without careful consideration of contextual variables can lead to unrealistic expectations and misallocation of resources.

9.0.4 Dynamic Nature of GenAI and Metrics Obsolescence

The GenAI landscape is characterized by rapid technological advancements, frequent model updates, and evolving best practices. ROI metrics established for an initial deployment might quickly become outdated as models improve, new features emerge, or the business context shifts. Continuous monitoring and adaptation of ROI measurement frameworks are therefore crucial, but this introduces overhead and complexity.

10 OPEN PROBLEMS AND RESEARCH GAPS

Despite the accelerating adoption of GenAI, several fundamental challenges remain largely unaddressed, representing significant open problems for researchers and practitioners alike.

10.0.1 *Standardized ROI Measurement Frameworks*

A lack of universally accepted, robust, and industry-agnostic frameworks for measuring GenAI ROI is a critical impediment. While various enterprises develop internal methodologies, these are rarely standardized, making cross-comparisons and aggregate industry analysis challenging. Future research should focus on developing a more generalized framework that can account for diverse use cases, industries, and organizational structures, integrating both quantitative and qualitative measures.

10.0.2 *Quantifying "Soft" and Strategic Benefits*

As discussed, translating intangible benefits such as enhanced creativity, accelerated market research, or improved employee engagement into quantifiable ROI figures remains an open challenge. Research is needed to develop reliable methodologies for:

- **Innovation ROI:** How GenAI contributes to novel product development or service offerings.
- **Employee Productivity Uplift:** Beyond simple task automation, measuring the cognitive load reduction and quality improvements from GenAI assistance.
- **Brand Perception and Trust:** The impact of GenAI-powered interactions on customer loyalty and brand reputation.

10.0.3 *Long-term vs. Short-term ROI Horizon*

Enterprises often prioritize short-term ROI to justify immediate investments. However, GenAI deployments typically involve significant upfront costs for infrastructure, data preparation, model training, and integration. The full strategic benefits and compounding returns may only materialize over a longer horizon. Balancing the imperative for short-term gains with strategic, long-term investments in GenAI capabilities, and developing ROI models that effectively capture both, is an ongoing problem. This includes accounting for ongoing maintenance, fine-tuning, and potential retraining costs associated with model drift or evolving business requirements.

10.0.4 *Scalability, Maintenance, and Governance Costs*

While pilot projects may show impressive ROI, scaling GenAI solutions across an entire enterprise introduces new cost dimensions, including robust MLOps pipelines, data governance frameworks, security protocols, and human oversight. Accurately forecasting and incorporating these ongoing operational expenses into ROI calculations is complex and often underestimated. The cost of managing model drift, ensuring data quality, and maintaining ethical compliance at scale are also significant considerations that need better quantification.

11 DATA QUALITY ISSUES

The efficacy and ROI of GenAI systems are profoundly dependent on the quality of the data they process and generate. Several data-related challenges can undermine ROI assessments.

11.0.1 *Availability of Granular and Clean Data*

Effective ROI measurement often requires granular data linking specific GenAI interactions or outputs to business outcomes. Many enterprises struggle with data silos, inconsistent data formats, and incomplete datasets, making it difficult to establish these connections. Data preparation, cleaning, and integration efforts can be substantial, consuming significant resources and potentially eroding initial ROI projections.

11.0.2 *Bias and Representativeness in Training Data*

Generative AI models are notorious for reflecting and amplifying biases present in their training data. If the data used to train or fine-tune enterprise GenAI models is unrepresentative or biased, the outputs can lead to skewed results, discriminatory outcomes, or inaccurate predictions, ultimately impacting business value and potentially leading to negative ROI through reputational damage or regulatory fines (Thukral et al., 2023). Identifying and mitigating such biases requires sophisticated data auditing and fairness-aware AI development practices.

11.0.3 *Data Privacy, Security, and Compliance*

The use of large datasets for GenAI, particularly those containing sensitive customer or proprietary enterprise information, raises significant data privacy and security concerns. Adherence to regulations such as GDPR, CCPA, and industry-specific compliance standards (e.g., HIPAA in healthcare) adds complexity and cost. Data breaches or misuse can severely impact reputation and incur substantial financial penalties, effectively rendering any positive GenAI ROI moot (Thukral et al., 2023). Secure data handling and anonymization techniques are critical but can limit data utility.

12 REVIEWER AUDIT: ANTICIPATED OBJECTIONS

Academic reviewers are likely to scrutinize the rigor and validity of any GenAI ROI claims. Several common objections can be anticipated:

12.0.1 *Selection Bias and Publication Bias*

Reviewers may question whether reported success stories represent a biased sample, focusing only on successful implementations while overlooking failures or projects with negative ROI. This "publication bias" can distort the overall understanding of GenAI's true enterprise value. A transparent discussion of failed projects or lessons learned from less successful deployments would strengthen the credibility of the research.

12.0.2 Lack of Robust Control Groups or Counterfactuals

A common critique in ROI studies is the absence of a true control group against which the GenAI intervention can be rigorously compared. Without a well-designed A/B test or a credible counterfactual, it is challenging to definitively attribute observed improvements solely to GenAI. Synthetic control methods or quasi-experimental designs (e.g., difference-in-differences) can help address this, but their application in complex enterprise settings can be difficult.

12.0.3 Over-attribution and Confounding Factors

Reviewers will likely challenge whether the observed ROI is truly incremental and directly attributable to GenAI, or if other confounding factors (e.g., new marketing strategies, economic upswings, organizational restructuring) have unduly influenced the results. This reiterates the need for robust causal inference methods, as discussed in Section 1.1.

12.0.4 Short-sightedness and Neglecting Total Cost of Ownership (TCO)

Focusing solely on immediate returns without accounting for the Total Cost of Ownership (TCO) over the entire lifecycle of a GenAI solution can be a significant oversight. This includes ongoing operational costs, maintenance, re-training, security, and the cost of managing associated risks. Reviewers will expect a holistic view of costs.

12.0.5 Reproducibility and Transparency

Given the proprietary nature of many enterprise GenAI deployments, detailed methodologies, data characteristics, and specific model configurations are often not publicly disclosed. This can lead to concerns about reproducibility and verifiability of ROI claims. Academic research benefits from transparency regarding methods, data sources (even if anonymized), and assumptions.

13 ETHICAL CONSIDERATIONS

Ethical implications are paramount in the deployment of GenAI and must be integrated into any comprehensive ROI assessment (Thukral et al., 2023). Failure to address these can lead to significant negative consequences, potentially negating any financial gains.

13.0.1 Algorithmic Bias and Fairness

Biased GenAI models can lead to unfair outcomes for specific customer segments or employee groups, resulting in reputational damage, legal challenges, and erosion of trust. For example, a GenAI model used for loan applications could inadvertently discriminate based on protected attributes if trained on historically biased data. Mitigating bias through fair AI principles, continuous monitoring, and impact assessments is crucial.

13.0.2 Transparency and Explainability

The "black box" nature of complex LLMs makes it difficult to understand how they arrive at specific outputs or recommendations. This lack of transparency can hinder trust, complicate error correction, and make it challenging to comply with explainability requirements in regulated

industries. Developing more interpretable GenAI systems and providing clear explanations for their decisions is an ongoing ethical imperative.

13.0.3 Job Displacement and Workforce Impact

The automation capabilities of GenAI raise concerns about potential job displacement. While GenAI can augment human capabilities and create new roles, its impact on the existing workforce requires careful consideration, including reskilling programs and ethical guidelines for deployment (Thukral et al., 2023). Ignoring these societal impacts can lead to negative public perception and regulatory backlash.

13.0.4 Misinformation, Hallucinations, and Safety

GenAI models, particularly LLMs, are prone to "hallucinating" false information or generating misleading content. In enterprise contexts, this can lead to incorrect business decisions, misinformation disseminated to customers, or unsafe recommendations. Robust validation, human oversight, and safety mechanisms are essential to mitigate these risks and maintain trust and accuracy.

13.0.5 Data Privacy and Security Breaches

As noted in data quality, the ethical handling of sensitive data is critical. Enterprises deploying GenAI must ensure stringent data governance, privacy-by-design principles, and robust security measures to prevent breaches and maintain customer trust. Ethical parameters and compliance are non-negotiable for sustainable GenAI ROI (Thukral et al., 2023).

In conclusion, while the promise of GenAI ROI is substantial, a thorough academic and practical understanding demands a rigorous engagement with its inherent limitations, open challenges, data considerations, anticipated criticisms, and profound ethical implications. A holistic and responsible approach is indispensable for realizing the true, sustainable value of GenAI in the enterprise.

REFERENCES

- [1] H. Soni, S. Z. Hassan, M. Deshapaga, M. Bansod, and R. N. Rajendran, "From tokens to tactics: Operationalizing generative ai in enterprise workflows," *Crossref*, 2026. [Online]. Available: <https://doi.org/10.2139/ssrn.7052339>
- [2] A. Zentner, "Agentic ai in the enterprise: A task-delegation taxonomy for autonomous agents, copilots, and human expertise," *Crossref*, 2026. [Online]. Available: <https://doi.org/10.2139/ssrn.7133258>
- [3] M. R., "Measuring business roi of generative ai adoption on azure cloud platforms," *EuropePMC Crossref*, 2026. [Online]. Available: <https://doi.org/10.21203/rs.3.rs-9033610/v1>
- [4] A. K. A. Kumar, "A causal roi framework for life sciences budget allocation, hcp targeting, and genai-driven personalization," *Crossref*, 2026. [Online]. Available: <https://doi.org/10.2139/ssrn.6374778>
- [5] A. Joshua, "Adoption depth and organizational-labor transformation in the ai era," *Crossref*, 2026. [Online]. Available: <https://doi.org/10.2139/ssrn.6323178>
- [6] S. Jamthe, "Generative ai for enterprise ai," *Crossref*, 2026. [Online]. Available: <https://doi.org/10.1201/9788743808145-14>
- [7] V. Thukral, L. Latvala, M. Swenson, and J. E. Horn, "Customer journey optimisation using large language models: Best practices and pitfalls in generative ai," *OpenAlex*, 2023. [Online]. Available: <https://openalex.org/W4400993506>

- [8] P. Arjunan, "Building enterprise genai infrastructure: A practical guide to storage requirements," *Crossref*, 2025. [Online]. Available: <https://doi.org/10.36948/ijfmr.2023.v05i01.19579>
- [9] T. Yiqing and S. Yang, "Layered architecture of enterprise ai agent infrastructure: an analytical framework based on cloud computing practice," *DOAJ*, 2026. [Online]. Available: <http://www.telecomsci.com/thesisDetails#10.11959/j.issn.1000-0801.DXKX260135>
- [10] H. Yang and H. Lee, "Leveraging generative ai for enterprise work: Integrating information adoption and task-technology fit model," *Crossref*, 2026. [Online]. Available: <https://doi.org/10.2139/ssrn.6065919>
- [11] P. Venkateela, "A framework for enterprise agentic ai: Multi-agent architectures and roi-driven adoption," *Crossref*, 2026. [Online]. Available: [https://doi.org/10.47363/jaicc/icadccs2026/2026\(5\)20](https://doi.org/10.47363/jaicc/icadccs2026/2026(5)20)
- [12] A. Rah, "Enterprise generative ai: A systematic review of security risks and governance," *Crossref*, 2026. [Online]. Available: <https://doi.org/10.2139/ssrn.6869661>
- [13] A. Bodnari and J. Travis, "Scaling enterprise ai in healthcare: the role of governance in risk mitigation frameworks," *OpenAlex Crossref DOAJ*, 2025. [Online]. Available: <https://openalex.org/W4410320336>
- [14] A. Rah, "Enterprise generative ai: A systematic review of security risks," *Crossref*, 2026. [Online]. Available: <https://doi.org/10.21203/rs.3.rs-9890887/v1>
- [15] L. Gyory, "Sla-driven orchestration of long-running multi-agent enterprise workflows," *Crossref*, 2025. [Online]. Available: <https://doi.org/10.2139/ssrn.5629770>
- [16] V. Tiwari, "Oracle ai: Agentic ai and enterprise applications transforming cloud erp, financial systems, and compliance through autonomous multi-agent architectures," *Crossref*, 2026. [Online]. Available: <https://doi.org/10.52710/cfs.1082>
- [17] C. Makridis, "Organizational transmission of ai: Managerial influence on generative ai adoption," *Crossref*, 2025. [Online]. Available: <https://doi.org/10.2139/ssrn.5401053>
- [18] G. N. C. Aluri and V. Manohar, "Generative ai integration patterns for enterprise workflow automation: A practitioner framework," *Crossref*, 2026. [Online]. Available: <https://doi.org/10.63282/3050-922x.ijeret-v6i3p121>
- [19] R. Modi, "Measuring business roi of generative ai adoption on azure cloud platforms," *OpenAlex EuropePMC*, 2026. [Online]. Available: <https://openalex.org/W7138188291>
- [20] AnkitaKapoor980, "Repository: Ankitakapoor980/genai-roi-powerbi-dashboard," *GitHub*, 2025. [Online]. Available: <https://github.com/AnkitaKapoor980/genai-roi-powerbi-dashboard>