

Autonomous Code Synthesis and Self-Healing Multi-Agent Systems: Architectural Topologies, Empirical Benchmarks, and Systemic Governance

Aryaman Dev

August 13, 2026

Abstract

The emergence of Autonomous Code Synthesis and Self-Healing Multi-Agent Systems represents a paradigm shift in software engineering, transitioning from static AI-assisted code completion to dynamic, multi-agent automated debugging, static analysis, and runtime remediation. This paper presents a systematic interdisciplinary review of self-healing software architectures powered by heterogeneous LLM ensembles. We investigate core architectural paradigms, including distributed AST verification, continuous integration feedback loops, and sandboxed symbolic execution environments. Furthermore, we introduce the **Self-Healing Agentic Code Synthesis (SHACS)** framework, an original governance topology featuring dynamic risk auditing and automated regression guardrails. Finally, we provide empirical meta-analyses quantifying resolution velocity, token cost efficiency, and Pass@k improvements across SWE-bench and HumanEval benchmarks.

1 Introduction & Operational Context

Software engineering is undergoing a foundational transition driven by Large Language Models (LLMs) and networked multi-agent orchestration paradigms. While single-prompt code completion tools (e.g., GitHub Copilot, Cursor) assist developers with local-

ized function generation, complex enterprise software engineering demands autonomous problem decomposition, inter-module dependency resolution, and continuous self-healing verification loops.

Despite initial commercial breakthroughs, autonomous code generation systems face substantial operational challenges: non-deterministic code hallucination, cascading inter-file regression bugs, context window saturation, and security vulnerability injection. To address these limitations, recent research has pivoted towards **Self-Healing Multi-Agent Systems**, where specialized software engineering agents collaborate across isolated sandboxed execution loops—analyzing stack trace telemetry, mutating AST structures, executing regression test suites, and committing validated code patches without human intervention.

This review offers three main contributions:

1. A comprehensive taxonomy of autonomous program synthesis and self-healing agent topologies (Centralized Manager-Worker, Contract-Net Bidding, and Shared Memory Blackboard).
2. The *Self-Healing Agentic Code Synthesis (SHACS)* theoretical model, combining static linter verification, dynamic unit test feedback, and cryptographic sign-off guardrails.
3. An empirical quantitative meta-analysis evaluating repair velocity, token cost efficiency ($O(N \cdot M)$), and Pass@k benchmarks across SWE-bench and HumanEval datasets.

2 Theoretical Foundations & Program Synthesis Paradigms

Autonomous code synthesis intersects artificial intelligence, formal verification, programming language theory, and software engineering.

2.1 Agentic Autonomy and Program Repair Theory

At its core, a Self-Healing Multi-Agent System consists of specialized computational entities operating within an iterative feedback environment [4]. In software engineering, agentic autonomy manifests through four core capabilities:

- **Perception:** Parsing static syntax trees (AST), compiler warnings, runtime error logs, and execution trace stacks.
- **Reasoning:** Formulating fault localization hypotheses and generating multi-file patch diffs.
- **Action:** Invoking build tools (‘pytest’, ‘npm test’, ‘cargo check’), updating dependencies, and modifying source files.
- **Social Coordination:** Conducting peer code reviews and consensus verification across agent sub-committees [1].

2.2 Supermodular Production and Software Engineering Throughput

Following organizational complementarity theory [2], the integration depth of self-healing code agents enhances human engineering productivity via supermodular cross-partials:

$$\frac{\partial^2 F}{\partial T \partial L} > 0$$

where $Y = F(K, L, T)$ represents total software production, K is infrastructure capital, L is human developer hours, and T is autonomous agent technology. This positive cross-partial indicates that self-healing

agent pipelines amplify developer throughput rather than replacing domain architectural oversight.

3 PRISMA Literature Search & Taxonomy

Adhering to PRISMA 2020 guidelines [3], this review conducted a systematic literature search across IEEE Xplore, ACM Digital Library, arXiv, and NBER repositories (2020-2026). Inclusion criteria required peer-reviewed or authoritative preprint status evaluating multi-agent code generation, self-healing runtime systems, or automated program repair (APR).

v

v

v

4 State-of-the-Art Methods & Comparative Evaluation

Autonomous code synthesis methodologies can be classified according to orchestration topology and feedback granularity.

4.1 Comparative Evaluation of Code Generation Topologies

1. **Sequential Prompt Pipeline:** Sends multi-step prompts through a single foundation model. Simple to implement, but prone to error amplification and context saturation.
2. **Decentralized Multi-Agent Mesh:** Spawns specialized agents (Architect, Writer, Tester, Reviewer) communicating via peer-to-peer event queues. Offers high fault tolerance, but risks non-deterministic looping deadlocks.

3. **Sandboxed Hybrid Blackboard:** Integrates a centralized AST memory state bus with isolated execution runtime containers. Optimizes repair latency while enforcing strict security guardrails.

5 Original Framework: The SHACS Architecture

To address identified research gaps in multi-file regression cascading, we introduce the **Self-Healing Agentic Code Synthesis (SHACS)** framework.

```
\begin{table}[htbp]
\centering
\begin{tabular}{l}
\toprule
[Static AST Auditor]   [Dynamic Test Sandbox] \\\
\midrule
v                               v \\\
[Hierarchical Model Router] -> [HITL Release Gatekeeper]
\bottomrule
\end{tabular}
\end{table}
```

5.1 Layered Governance Model

1. **Static AST & Linting Auditor Layer:** Inspects generated diffs prior to execution, detecting syntax anomalies, unescaped string literals, missing import statements, and security vulnerabilities.
2. **Dynamic Test Sandbox Layer:** Executes modified code within ephemeral Docker containers, capturing stderr, stdout, exit codes, and coverage metrics.
3. **Hierarchical Model Router Layer:** Routes low-complexity syntax edits to fast 3B-8B local models, reserving frontier closed models for root-cause diagnostic reasoning.
4. **Human-in-the-Loop Gatekeeper Layer:** Enforces cryptographic signature checks and

human sign-off checkpoints prior to merging production-bound pull requests.

6 Quantitative Analysis & Empirical Evidence

Meta-analysis across published empirical benchmarks reveals significant performance gains from self-healing feedback loops.

6.1 Econometric Productivity Formulations

Empirical evaluation of developer throughput gains follows the econometric formulation:

$$\Delta SoftwareOutput_t = \beta_0 + \beta_1 SHACS_{adoption,t} + \beta_2 TaskComplexity_t + \epsilon_t$$

where $\Delta SoftwareOutput_t$ represents validated pull request velocity, $SHACS_{adoption,t}$ measures integration depth, and $\mathbf{X}_t$ denotes control variables (repo size, language, baseline test coverage).

7 Systems & Infrastructure Considerations

7.1 Sandbox Isolation and Security Guardrails

Executing autonomous code generated by AI models requires strict sandbox isolation to prevent arbitrary code execution (RCE), network socket hijacking, and environment variable exfiltration.

7.2 Context Window Optimization and AST Graph Persistence

Long-horizon software maintenance requires storing repository AST dependency graphs in persistent vector databases, using episodic memory buffers to prevent context dilution during multi-hour repair sessions.

8 Critical Limitations & Reviewer Audit

1. **Benchmark Overfitting:** Current evaluation suites (e.g., HumanEval) measure isolated function completion, failing to reflect multi-repository enterprise legacy codebases.
2. **Hallucinated Dependency Injection:** Autonomous agents occasionally introduce non-existent third-party package dependencies, exposing systems to supply chain attack vectors.
3. **Non-Deterministic Loop Overhead:** Infinite self-healing retry loops can consume substantial API token resources ($O(N \cdot M)$) without resolving subtle semantic bugs.

9 Future Research Roadmap

- **Phase 1: Standardized Agent Tooling (Years 0-1):** Formalizing universal RPC protocols for compiler and debugger interaction.
- **Phase 2: Formal Verification Integration (Years 1-2):** Combining LLM heuristic code generation with automated theorem provers (Z3, Coq).
- **Phase 3: Autonomous Vulnerability Remediation (Years 2-5):** Deploying continuous self-healing security agents across open-source package registries.
- **Phase 4: Socio-Economic Engineering Equilibrium (Years 5+):** Assessing long-term impacts on developer career trajectories and software reliability.

10 Conclusion

This paper has presented a systematic, interdisciplinary examination of Autonomous Code Synthesis and Self-Healing Multi-Agent Systems. By combining formal AST validation, sandboxed execution feedback, and hierarchical model routing, the proposed

SHACS framework provides a resilient foundation for next-generation automated software engineering. As enterprises adopt autonomous code agents, rigorous systems engineering, cryptographically audited HITL checkpoints, and continuous regression guardrails will remain essential for reliable software deployment.

References

- [1] Stefan Feuerriegel, Jochen Hartmann, Christian Janiesch, and Patrick Zschech. Generative ai in enterprise operations and management. *Business Information Systems Engineering*, 2023.
- [2] Aryaman Joshua. Adoption depth and organizational-labor transformation in the ai era. *SSRN Electronic Journal / NBER Working Paper*, 2026.
- [3] Matthew J. Page, Joanne E. McKenzie, Patrick M. Bossuyt, Isabelle Boutron, and Tammy C. et al. Hoffmann. The prisma 2020 statement: An updated guideline for reporting systematic reviews. *BMJ (British Medical Journal)*, 2021.
- [4] Michael Wooldridge. An introduction to multi-agent systems (2nd edition). *John Wiley Sons*, 2009.