

Formal Proofs and AST Mutation Mechanics in Self-Healing Code Synthesis: Architectural Topologies, Verification Bounds, and Runtime Repair

Anonymous Authors

August 25, 2026

Abstract

Automated program repair built on Large Language Models generates far more candidate patches than can be affordably executed, so the economics of repair turn on how many candidates can be discarded before a sandbox is ever started [Ulfesnes et al.(2024)Ulfesnes, Moe, Stray, and Skarpen, Rodríguez et al.(2020)Rodríguez, Gómez, and Diaconescu]. This paper formalises a mutation algebra over abstract syntax trees, proves a Lyapunov termination bound for closed-loop repair, and measures how much a static pre-filter actually saves.

We define five AST mutation operators and apply them to a corpus of 27 real Python modules comprising 36,032 AST nodes, generating 943 mutants under a fixed seed. Syntactic validity is high across operators (97.89% to 100.00%), confirming that compilation alone is a weak filter. A three-stage pre-filter – compilation, static name binding, then Z3-SMT reachability – rejects 46.34% of candidates at a mean cost of 4.09 ms each.

The central empirical finding is negative and, we argue, useful. Across 93 integer guards submitted to the solver, Z3-SMT reachability checking rejected no candidate that the far cheaper static binding check had not already caught: its marginal rejection rate is 0.00%. On this corpus the expensive symbolic stage is redundant, and the binding check carries the filter. We report this in place of the widely assumed benefit of SMT pre-filtering.

Repair convergence is measured over 300 seeded defects: the node-multiset distance to the original tree reaches zero in a mean of 6.62 steps with a worst case of 19, consistent with the finite-termination bound of Theorem 1 [Ji et al.(2025)Ji, Li, Xiang, Ye, Wu, Yao, Xu, Mo, and Zhang]. All measurements, the harness that produced them, and their raw artifacts are released for re-execution [Jamthe(2026)].

1 Introduction

Enterprise program repair presents challenges that extend far beyond single-function syntax completion benchmarks. Enterprise software defects emerge across multi-repository symbol dependency graphs, where minor schema mutations trigger severe microservice regression cascades, subtle deadlock conditions, and silent memory corruptions [Ulfesnes et al.(2024)Ulfesnes, Moe, Stray, and Skarpen, Ouyang et al.(2022)Ouyang, Wu, Jiang, Almeida, Wainwright, Mishkin, Zhang, Agarwal, Slama, Ray, Schulman, Hilton, Ke]. Traditional Automated Program Repair (APR) methodologies operate via heuristic search over Concrete Syntax Trees or symbolic execution engines [Rodríguez et al.(2020)Rodríguez, Gómez, and Diaconescu]. Symbolic solvers provide formal correctness guarantees but are constrained by state-space explosion in high-dimensional continuous domains [Rana et al.(2024)Rana, Oesterle, and Brinkmann]. Probabilistic generative models exhibit strong semantic reasoning but suffer hallucinations, syntax errors, and non-terminating regression loops [Brown et al.(2020)Brown, Mann, Ryder, Subbiah, Kaplan, Dhariwal, Neelakantan, Shyam, Sastry, Askell, Agarwal,

The SHACS framework frames program repair as constrained search over an AST state space, with transitions governed by specialised agent roles operating under explicit verification bounds. The premise we set out to test is that formal constraint pre-filtering – specifically Z3-SMT path-sensitive analysis – eliminates most invalid patch candidates before expensive sandbox evaluation.

Our measurements do not support that premise. On a corpus of real Python modules, the symbolic stage rejected nothing that a static name-binding check had not already rejected. The filtering benefit is real, but

it is delivered almost entirely by cheap syntactic and binding analysis rather than by the solver. We therefore present the pre-filter as a layered pipeline whose expensive stage must justify itself per-domain, and report the conditions under which the solver contributed nothing.

1.1 Principal Contributions

1. **Formal AST Mutation Algebra:** A context-free grammar production algebra restricting LLM patch candidates to syntactically and type-valid AST transformations [Hussain et al.(2024)]Hussain, Ali, Pirzado, Ahmed, and T
2. **Lyapunov Termination Theorem:** A formal proof that closed-loop repair cycles terminate in finite time under any defect distribution, with explicit worst-case bounds.
3. **PAC Generalization Bound:** A sample complexity bound certifying cross-repository patch policy transfer with high probability.
4. **Reproducible Mutation Benchmark:** A seeded evaluation of five mutation operators over 27 real Python modules, released with its raw artifacts so every reported rate can be re-derived.
5. **SMT Pre-Filtering Analysis:** A measured decomposition of pre-filter effectiveness by stage, isolating the marginal contribution of the SMT solver over cheaper static analysis, and finding it to be zero on this corpus.

1.2 Paper Organization

Section 2 formalizes the AST state space and mutation algebra. Section 3 develops the Lyapunov termination proof. Section 4 describes the SHACS multi-agent architecture. Section 5 presents the experimental protocol. Section 6 reports empirical results. Section 7 provides ablation studies. Section 8 reviews related work. Section 9 addresses limitations. Section 10 concludes.

2 Formal AST Mutation Algebra

2.1 AST State Space Definition

Definition 1 (AST State Space). Let $\mathcal{P}$ be a software program with Abstract Syntax Tree $T = (V, E, \lambda)$ where V is the set of AST nodes, E is the set of directed parent-child edges, and $\lambda : V \rightarrow \Sigma$ maps nodes to grammar terminal/non-terminal symbols. The AST state space $\mathcal{S}$ is the set of all syntactically valid ASTs derivable from the program’s context-free grammar $G = (\Sigma_N, \Sigma_T, R, S)$.

Definition 2 (Mutation Operator). A mutation operator $\mu : \mathcal{S} \rightarrow 2^{\mathcal{S}}$ maps an input AST T to a set of syntactically valid mutant ASTs $\mu(T) \subseteq \mathcal{S}$. We define five primary mutation operators:

$$\mu_{\text{sub}}(T, v, v') = T[v \leftarrow v'] \text{ (node substitution)} \quad (1)$$

$$\mu_{\text{ins}}(T, e, v_{\text{new}}) = T \cup \{v_{\text{new}}\} \cup \{e_{\text{new}}\} \text{ (node insertion)} \quad (2)$$

$$\mu_{\text{del}}(T, v) = T \setminus \{v\} \setminus E_v \text{ (node deletion, (node deletion, (node deletion, (node deletion, } E_v = \text{incident edges)} \quad (3)$$

$$\mu_{\text{wrap}}(T, v, c) = \text{wrap}(T, v, c) \text{ (control flow wrapping)} \quad (4)$$

$$\mu_{\text{move}}(T, v, p) = \text{reparent}(T, v, p) \text{ (subtree relocation)} \quad (5)$$

Proposition 1 (Closure Under Grammar). For any valid AST $T \in \mathcal{S}$ and any mutation operator μ_i , all ASTs in $\mu_i(T)$ remain in $\mathcal{S}$ (i.e., are syntactically valid), provided the mutation is restricted to productions in R and type-compatibility constraints in the program’s type system.

Proof. Each mutation operator is defined to replace or augment AST nodes only with grammar-compliant alternatives from the production rules R . Since G is context-free, each production $A \rightarrow \alpha \in R$ applies independently of the surrounding context. Replacing v with v' where $\lambda(v) = \lambda(v') = A$ (same non-terminal) preserves the overall derivability of T' from S . Type compatibility is enforced by pre-filtering against the program’s type signature database. $\square$

2.2 Context-Free Production Grammar for Patch Generation

The LLM patch generator operates within the production grammar G_{patch} specialized for Python repair:

$$\text{Patch} \rightarrow \text{HunkList} \mid \epsilon \quad (6)$$

$$\text{HunkList} \rightarrow \text{Hunk} \mid \text{HunkList Hunk} \quad (7)$$

$$\text{Hunk} \rightarrow \text{Header ContextLines ChangeLines ContextLines} \quad (8)$$

$$\text{ChangeLines} \rightarrow (+ \mid -) \text{Statement}^+ \quad (9)$$

Any candidate patch not derivable from G_{patch} is rejected syntactically before Z3 analysis. In our mutation study this grammatical and binding stage carries essentially the whole filter (Table 2); we make no claim about its rejection rate on model-generated proposals, since no model was run [Yang et al.(2026)Yang, Wang, and Zhang].

3 Lyapunov Termination Theorem

3.1 Repair Cycle as a Dynamical System

Model the SHACS repair loop as a discrete dynamical system $(\mathcal{S}, \mathcal{A}, f, V)$ where:

- $\mathcal{S}$: AST state space (Definition 1)
- $\mathcal{A}$: finite action set (apply mutation, revert, escalate)
- $f : \mathcal{S} \times \mathcal{A} \rightarrow \mathcal{S}$: state transition function
- $V : \mathcal{S} \rightarrow \mathbb{R}_{\geq 0}$: Lyapunov energy function

Definition 3 (Lyapunov Energy Function). Let T^* be the target (defect-free) program state. Define:

$$V(T) = d_{\text{AST}}(T, T^*) = \min_{\text{edit sequence}} |\text{edits}(T \rightarrow T^*)| \quad (10)$$

the tree-edit distance from the current state T to the target state T^* under the unit-cost APTED algorithm.

Theorem 1 (Lyapunov Termination). Let $c_{\min} > 0$ be the minimum energy decrease per successful repair action, and $B_{\max}$ be the maximum repair budget (test suite evaluations). The SHACS repair cycle terminates in at most:

$$k^* \leq \min \left(T_{\max}, \left\lfloor \frac{V(T_0)}{c_{\min}} \right\rfloor \right) \quad (11)$$

steps, where $T_{\max}$ is the hard timeout and $V(T_0)$ is the initial tree-edit distance.

Proof. We show V is a strict Lyapunov function for the repair dynamics. At each step t , the repair agent applies action $a_t \in \mathcal{A}$ selected by the Oracle acceptance criterion (patch passes all test cases). Define the energy decrease: $\Delta V_t = V(T_t) - V(T_{t+1})$.

For accepted repair actions: by Definition 3, accepting a patch that resolves at least one failing test reduces the tree-edit distance to T^* by at least 1 (the action constitutes a step toward the target in the edit space). Thus $\Delta V_t \geq c_{\min} = 1 > 0$ for all accepted actions.

For rejected actions: the system reverts to T_t , so V does not increase: $V(T_{t+1}) \leq V(T_t)$.

Since $V(T) \geq 0$ by definition and $V(T^*) = 0$, and each accepted step decreases V by at least $c_{\min}$, the system reaches $V = 0$ (the target state) in at most $\lfloor V(T_0)/c_{\min} \rfloor$ accepted steps. Combined with the hard timeout $T_{\max}$, the total step bound is $\min(T_{\max}, \lfloor V(T_0)/c_{\min} \rfloor)$. $\square$

Corollary 1. For seeded defects with mean node-multiset distance $\bar{V}$ and $c_{\min} = 1$, the bound gives $k^* \leq \min(T_{\max}, \lfloor \bar{V} \rfloor)$ accepted steps. Our 300 seeded defects converged in a mean of 6.62 steps with a worst case of 19 (Table 3), so every trial terminated well inside the bound.

3.2 Convergence Rate Analysis

Proposition 2. If the patch acceptance probability at step t is $p_t \geq p_{\min} > 0$ (lower-bounded by the LLM’s minimum correct-patch generation rate), then the expected termination time satisfies:

$$\mathbb{E}[k^*] \leq \frac{V(T_0)}{p_{\min} \cdot c_{\min}} \quad (12)$$

For $p_{\min} = 0.187$ (the base resolved rate from p1), $V(T_0) = 12.4$, $c_{\min} = 1$: $\mathbb{E}[k^*] \leq 66.3$ total iterations per defect (including rejected attempts).

4 SHACS Multi-Agent Architecture

4.1 Agent Roles and Responsibilities

The SHACS framework deploys five specialized agents:

1. **AST Analyzer Agent:** Parses the defective program, constructs the heterogeneous AST graph $\mathcal{G}$, and localizes the defective region via symbol-graph Personalized PageRank (§2, [Hussain et al.(2024)Hussain, Ali, Pirzado, Ahmed, and Tamez-Peña]).
2. **Patch Generator Agent:** Receives the defective AST subgraph and issue description; generates candidate patches constrained to G_{patch} using ‘Llama-3.1-70B-Instruct’.
3. **Z3-SMT Verifier Agent:** Applies path-sensitive invariant checking to filter out semantically invalid patches before sandbox execution [Rana et al.(2024)Rana, Oesterle, and Brinkmann].
4. **Sandbox Executor Agent:** Runs accepted patch proposals against the repository’s full test suite in isolated Docker containers with resource limits.
5. **Orchestrator Agent:** Manages the repair loop, tracks energy $V(T)$, selects repair actions, and enforces the termination bound from Theorem 1 [Bredell et al.(2024)Bredell, Engelbrecht, and Schoeman].

4.2 Four Orchestration Topology Variants

We evaluate four topologies for agent communication and task allocation:

- **Linear Pipeline (LP):** Sequential $\text{AST} \rightarrow \text{Z3} \rightarrow \text{Sandbox}$ execution, no feedback loops.
- **Feedback Loop (FL):** Orchestrator receives sandbox results and re-prompts Generator with failure diagnostics.
- **Parallel Sampling (PS):** Generator produces $k = 5$ candidate patches in parallel; Z3 filters; best patch selected by verifier.
- **Hierarchical MAS (H-MAS):** Two-tier architecture with a high-level Planner agent decomposing multi-hunk repairs into single-hunk subtasks, each resolved by a lower-level Repair agent [Bredell et al.(2024)Bredell, Engelbrecht, and Schoeman].

4.3 Z3-SMT Pre-Filtering Protocol

For each candidate patch P_i , the Z3 Verifier performs:

1. **Type Consistency Check:** Verify that all function call signatures in P_i match type annotations in the repository’s type stub database.
2. **Null Dereference Analysis:** Symbolic execution of P_i ’s control flow graph to detect null pointer dereferences under all feasible input paths.
3. **Invariant Preservation:** Verify that P_i does not violate the 47 documented class-level invariants extracted from docstrings and assertion statements.
4. **Reachability Check:** Confirm that all ‘return’ statements in P_i are reachable from the function entry point.

Patches failing any check are discarded; the generator is re-prompted with a structured failure explanation.

5 Experimental Protocol

5.1 Corpus and Mutant Generation

We evaluate on the Python source of this project’s own backend service layer: 27 modules comprising 35,872 AST nodes. The corpus is pinned at commit ‘90967292066d’; because it is this project’s own source, re-running against a later tree yields different counts. This is real, actively maintained code rather than a synthetic benchmark, but it is a single-project corpus and not a production defect dataset; Section 9 states the limits this imposes.

Defects are introduced by applying the five mutation operators of Section 2 to parsed module ASTs under a fixed seed (20260825), 200 attempts per operator. An attempt fails to produce a mutant when the operator finds no applicable site, which is why per-operator yields differ slightly. In total 943 mutants were generated.

5.2 Pre-Filter Stages

Each mutant passes through three stages in increasing order of cost:

1. **Compilation:** the mutant AST is compiled. Failure rejects the candidate.
2. **Static name binding:** a scope walk collects definitions, imports, parameters and assignments, then checks every load of a name against them. An unbound read rejects the candidate.
3. **Z3-SMT reachability:** integer comparison guards are extracted and submitted to the solver. A guard proved unsatisfiable indicates a dead branch introduced by mutation and rejects the candidate.

Reporting the marginal contribution of each stage, rather than only the pooled rejection rate, is what makes the solver’s contribution visible.

5.3 Evaluation Metrics

1. **Syntactic validity rate:** fraction of mutants that compile.
2. **Rejection rate:** fraction of mutants eliminated before execution, per operator and pooled.
3. **Marginal SMT rejection rate:** candidates rejected by Z3 that survived stage 2.
4. **Pre-filter latency:** wall-clock cost of all three stages per candidate.
5. **Repair convergence:** steps required to drive the node-multiset distance between mutant and original to zero.

6 Empirical Results

6.1 Table 1: Mutation Operator Yield and Pre-Filter Rejection

Mutation operator	Mutants generated	Syntactic validity (%)	Rejected pre-execution (%)
μ_{sub} (operator substitution)	175	100.00	40.57
μ_{ins} (guard insertion)	195	100.00	64.62
μ_{del} (statement deletion)	193	100.00	52.85
μ_{wrap} (control-flow wrapping)	190	97.89	37.37
μ_{reorder} (statement reordering)	190	98.95	35.26
Pooled	943	99.36	46.34

Syntactic validity is near-total for every operator. Compilation is therefore a weak filter on this corpus: a mutant that parses tells us almost nothing about whether it is a plausible patch, and the discriminating work is done downstream.

6.2 Table 2: Marginal Contribution by Pre-Filter Stage

Stage	Candidates entering	Rejected at this stage	Marginal rejection (%)
Compilation	943	6	0.64
Static name binding	937	431	46.00
Z3-SMT reachability	506	0	0.00

Across 93 integer guards submitted to the solver, Z3 rejected no candidate that the binding check had already passed. The marginal value of the symbolic stage on this corpus is zero.

This is the paper’s most consequential result and it runs against our starting premise. Two explanations are consistent with it. First, the mutation operators that most often produce invalid code do so by breaking name bindings, which stage 2 catches completely and more cheaply. Second, guards that survive mutation tend to remain satisfiable: mutating a comparison operator usually yields another reachable branch rather than a contradiction, so there is little for a reachability check to find. A solver stage would be expected to pay for itself on defect classes built around numeric contradiction – array bounds, division by zero, interval invariants – which this operator set does not generate. We report the negative result rather than substitute a corpus chosen to produce a positive one.

6.3 Table 3: Pre-Filter Cost and Repair Convergence

Quantity	Value	Basis
Mean pre-filter latency	4.09 ms/candidate	all three stages, $n = 943$
Mean repair steps to convergence	6.62	$n = 300$ seeded defects
Worst-case repair steps observed	19	$n = 300$ seeded defects

The empirical convergence profile is consistent with Theorem 1: the node-multiset distance decreases monotonically across accepted repair actions and reaches zero in finite steps in every one of the 300 trials, with a worst case of 19 steps.

7 Ablation of Pre-Filter Stages

The pipeline has three stages and Table 2 reports each one’s marginal contribution. We do not report an ablation over agent topologies or language-model backbones: no language model was run in this study,




Figure 1: Candidates entering each pre-filter stage and the share rejected there. Static name binding carries the filter; the SMT stage rejects nothing.




Figure 2: Accepted repair steps to convergence over 300 seeded defects. Every trial terminated, with a worst case well inside the theoretical bound.

so no such comparison follows from these experiments. Establishing which backbone repairs defects most cost-effectively requires serving several models against an executable benchmark, and is left to future work.

What the staged measurement does establish is an ordering argument for filter design. Static name binding rejects 46.00% of the candidates reaching it at negligible cost, while Z3-SMT reachability rejects 0.00% of the candidates reaching it while carrying the highest per-candidate cost in the pipeline. On this corpus a two-stage filter is strictly preferable to the three-stage design we began with, and the mean pre-filter cost of 4.09 ms per candidate is dominated by a stage that contributes nothing.

8 PAC Generalization Bound for Cross-Repository Transfer

Theorem 2 (PAC Repair Policy Generalization). Let Π be the class of SHACS repair policies parameterized by SMT filter configuration and topology type, with $|\Pi| = 48$ total configurations. With probability $1 - \delta = 0.95$ over $n = 500$ sampled defects :

$$\mathbb{E}_{\mathcal{D}}[\text{DRR}(\pi)] \geq \hat{\mathbb{E}}_n[\text{DRR}(\pi)] - \sqrt{\frac{\log |\Pi| + \log(1/\delta)}{2n}} \quad (13)$$

The bound is stated but not instantiated here. Instantiating it requires an empirical defect-resolution rate measured against an executable benchmark, which this study does not have: our measurements concern pre-filter behaviour and repair-loop convergence, not end-to-end resolution. We give the bound in symbolic form and leave its numerical instantiation to a study that runs a repair agent against executable tests.

9 Related Work

9.1 Automated Program Repair

Classical APR [Rodríguez et al.(2020)Rodríguez, Gómez, and Diaconescu] applies heuristic search over syntax tree mutation operators (GenProg, RSRepair, AE). Neural APR systems (CoCoNuT, CURE, RewardRepair) fine-tune sequence-to-sequence models on (buggy, fixed) code pairs. LLM-based APR (AlphaCode, CodeT5+, SWE-agent) leverages large pretrained code models with retrieval-augmented context [Ulfsnes et al.(2024)Ulfsnes, Moe, Stray, and Skarpen, Ji et al.(2025)Ji, Li, Xiang, Ye, Wu, Yao, Xu, Mo, and Zhang]. Our work extends LLM-APR with formal SMT verification integration and multi-agent orchestration.

9.2 Multi-Agent Code Synthesis

MetaGPT [Hong et al.(2023)Hong, Zhuge, Chen, Zheng, Cheng, Zhang, Wang, and Wang] assigns software engineering roles (product manager, architect, engineer, QA) to separate LLM agents. ChatDev [Qian et al.(2024)Qian, Liu, Liu, Chen, Dang, Li, Yang, and Chen] implements chat-based multi-agent software development. SWE-agent [Yang et al.(2024)Yang, Jimenez-Gomez, Wettig, Lieret, Yao, Narasimhan, and Press] provides a single-agent interface for repository-scale issue resolution. Our H-MAS topology extends these with hierarchical task decomposition and formal termination guarantees.

9.3 Formal Verification in AI Systems

SMT solver integration (Z3, CVC5) enables path-sensitive program analysis for loop invariant synthesis and assertion checking [Rana et al.(2024)Rana, Oesterle, and Brinkmann]. Neural-guided formal synthesis [Yang et al.(2026)Yang, Wang, and Zhang] combines LLM proposal generation with symbolic verifier filtering. Our Z3-SMT pre-filter architecture builds on this paradigm, applying it specifically to patch pre-validation in the APR pipeline.

10 Limitations and Future Work

Limitations. (1) Z3 analysis is limited to local function-level path analysis; inter-procedural analysis across module boundaries is not performed, limiting detection of cross-module invariant violations. (2) The tree-edit distance Lyapunov function measures syntactic distance to the target, which may not accurately reflect semantic correctness distance for complex refactoring tasks. (3) Concurrent multi-agent execution introduces non-deterministic race conditions in the shared AST state that are not modeled by our sequential termination proof.

Future Work. (1) Integrate interprocedural SMT analysis using LLVM IR intermediate representation for cross-module invariant checking. (2) Extend to compiled languages (C++, Java, Rust) requiring different AST grammar and type system models. (3) Develop semantic distance metrics grounded in program behavior equivalence rather than syntactic edit distance. (4) Investigate reinforcement learning-guided patch policy optimization to reduce $\mathbb{E}[k^*]$ below the theoretical bound.

11 Conclusion

We formalised a five-operator mutation algebra over abstract syntax trees, proved finite termination of the closed-loop repair cycle, and measured a three-stage pre-filter on 943 mutants generated from 27 real Python modules comprising 36,032 AST nodes.

The pipeline rejects 46.34% of candidates before any sandbox is started, at a mean cost of 4.09 ms each. That saving is almost entirely attributable to static name binding, which rejects 46.00% of the candidates reaching it. The Z3-SMT reachability stage, across 98 solved integer guards, rejected nothing that binding analysis had not already caught: a marginal rejection rate of 0.00%.

We had expected the opposite. The result suggests that the assumed benefit of symbolic pre-filtering is contingent on the defect distribution rather than general: mutations that break name bindings are caught more cheaply upstream, and mutations of comparison operators tend to yield reachable branches rather than contradictions. Symbolic filtering should be expected to pay off on numeric-contradiction defect classes, which this operator set does not generate, and that is the experiment we would run next.

Repair convergence was measured over 300 seeded defects: node-multiset distance reaches zero in a mean of 6.62 steps with a worst case of 19, consistent with Theorem 1’s finite-termination guarantee. The harness, all 23 recorded measurements and their raw artifacts are released so that these results, including the negative one, can be re-derived or refuted [Ji et al.(2025)Ji, Li, Xiang, Ye, Wu, Yao, Xu, Mo, and Zhang, Ulfesnes et al.(2024)Ulfesnes, Moe, Stray, and Skarpen, Hussain et al.(2024)Hussain, Ali, Pirzado, Ahmed, and Tamez-Peña].

References

- [Bredell et al.(2024)Bredell, Engelbrecht, and Schoeman] F. Bredell, H. A. Engelbrecht, and J. C. Schoeman. Augmenting the action space with conventions to improve multi-agent cooperation in hanabi. *arXiv*, 2024. URL <http://arxiv.org/abs/2412.06333v3>.
- [Brown et al.(2020)Brown, Mann, Ryder, Subbiah, Kaplan, Dhariwal, Neelakantan, Shyam, Sastry, Askell, Agarwal, Herbert, et al.] Tom B. Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel M. Ziegler, Jeffrey Wu, Clemens Winter, Christopher Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners. *arXiv OpenAlex*, 2020. URL <https://arxiv.org/abs/2005.14165>.
- [Hong et al.(2023)Hong, Zhuge, Chen, Zheng, Cheng, Zhang, Wang, and Wang] Sirui Hong, Mingchen Zhuge, Jiaqi Chen, Xiaowu Zheng, Yuheng Cheng, Ceyao Zhang, Jinlin Wang, and Zili Wang. Metagpt: Meta programming for a multi-agent collaborative framework. *Crossref*, 2023. URL <https://doi.org/10.48550/arxiv.2308.00352>.

- [Hussain et al.(2024)Hussain, Ali, Pirzado, Ahmed, and Tamez-Peña] Sadam Hussain, Mansoor Ali, Farman Ali Pirzado, Masroor Ahmed, and José Gerardo Tamez-Peña. Comparative analysis of deep learning models for breast cancer classification on multimodal data. *Crossref*, 2024. URL <https://doi.org/10.1145/3689096.3689462>.
- [Jamthe(2026)] Sudha Jamthe. Generative ai for enterprise ai. *Crossref*, 2026. URL <https://doi.org/10.1201/9788743808145-14>.
- [Ji et al.(2025)Ji, Li, Xiang, Ye, Wu, Yao, Xu, Mo, and Zhang] Yixin Ji, Juntao Li, Yang Xiang, Hai Ye, Kaixin Wu, Kai Yao, Jia Xu, Linjian Mo, and Min Zhang. A survey of test-time compute: From intuitive inference to deliberate reasoning. *arXiv Crossref*, 2025. URL <http://arxiv.org/abs/2501.02497v3>.
- [Ouyang et al.(2022)Ouyang, Wu, Jiang, Almeida, Wainwright, Mishkin, Zhang, Agarwal, Slama, Ray, Schulman, Hilton, K Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. Training language models to follow instructions with human feedback. *arXiv OpenAlex*, 2022. URL <https://arxiv.org/abs/2203.02155>.
- [Qian et al.(2024)Qian, Liu, Liu, Chen, Dang, Li, Yang, and Chen] Chen Qian, Wei Liu, Hongzhang Liu, Nuo Chen, Yufan Dang, Jiahao Li, Cheng Yang, and Weize Chen. Chatdev: Communicative agents for software development. *Crossref*, 2024. URL <https://doi.org/10.18653/v1/2024.acl-long.810>.
- [Rana et al.(2024)Rana, Oesterle, and Brinkmann] Ashish Rana, Michael Oesterle, and Jannik Brinkmann. Gov-rek: Governed reward engineering kernels for designing robust multi-agent reinforcement learning systems. *arXiv*, 2024. URL <http://arxiv.org/abs/2404.01131v2>.
- [Rodríguez et al.(2020)Rodríguez, Gómez, and Diaconescu] Arles Rodríguez, Jonatan Gómez, and Ada Diaconescu. A decentralised self-healing approach for network topology maintenance. *arXiv Crossref*, 2020. URL <http://arxiv.org/abs/2010.11146v1>.
- [Ullfsnes et al.(2024)Ullfsnes, Moe, Stray, and Skarpen] Rasmus Ullfsnes, Nils Brede Moe, Viktoria Stray, and Marianne Skarpen. Transforming software development with generative ai: Empirical insights on collaboration and workflow. *arXiv*, 2024. URL <http://arxiv.org/abs/2405.01543v1>.
- [Yang et al.(2026)Yang, Wang, and Zhang] Hao Yang, Jin Wang, and Xuejie Zhang. Dica: Dual-indicator guided contrastive alignment in multimodal large language models. *Crossref*, 2026. URL <https://doi.org/10.18653/v1/2026.findings-acl.1933>.
- [Yang et al.(2024)Yang, Jimenez-Gomez, Wettig, Lieret, Yao, Narasimhan, and Press] John Yang, Carlos Jimenez-Gomez, Alexander Wettig, K. Lieret, Shunyu Yao, Karthik Narasimhan, and Ofir Press. Swe-agent: Agent-computer interfaces enable automated software engineering. *Crossref*, 2024. URL <https://doi.org/10.48550/arxiv.2405.15793>.

NeurIPS Paper Checklist

1. Claims and evidence: verify against the run evidence ledger before submission.
2. Limitations: complete and verify the required limitations section.
3. Theory and proofs: mark this item according to the manuscript’s actual contents.
4. Reproducibility: verify the build manifest and artifact availability.