

Cryptographic Attestation and Hardware-Backed Gatekeeping for Microservice Meshes: Hardware Security Modules, Ephemeral Sandboxes, and Ed25519 Audits

Anonymous Authors

August 20, 2026

Abstract

As enterprise software architectures transition toward distributed multi-agent artificial intelligence (AI) systems, traditional microservice boundary security model becomes vulnerable to identity impersonation, prompt injection, and lateral privilege escalation. This paper presents a novel framework for Cryptographic Attestation and Hardware-Backed Gatekeeping tailored for enterprise multi-agent microservice meshes. We introduce an architecture integrating Hardware Security Modules (HSMs), ephemeral WebAssembly (Wasm) execution sandboxes, and Ed25519 digital signature audits to enforce zero-trust identity verification across inter-agent communications. Our empirical evaluation across a 500-agent mesh demonstrates that hardware-backed attestation introduces less than 4.2 ms of cryptographic overhead per transaction while reducing unauthorized execution attempts to zero ($p < 0.001$). [[crossref_10.1109_access.2026.

Enterprise deployment of autonomous multi-agent systems introduces significant challenges to system governance, identity verification, and runtime security [1]. Unlike static microservices with deterministic API call graphs, multi-agent frameworks exhibit emergent communication patterns, dynamic agent spawning, and variable privilege boundaries.

Traditional OAuth2 and TLS mutual authentication mechanisms fall short when protecting autonomous agents operating across untrusted or semi-trusted execution environments. Without cryptographic proof of agent binary state and memory integrity, an adversary executing prompt injection can hijack an agent process and issue unauthorized transactions across the service mesh.

To address these vulnerabilities, we introduce a hardware-enforced security framework incorporating three core mechanisms:

1. *Root-of-Trust Hardware Attestation: Utilizing HSMs and Trusted Execution Environments (TEEs) to sign agent identity certificates upon initialization.*
2. *Ephemeral Sandboxing: Spawning agent tasks within isolated, webassembly-based micro-runtimes with bounded VRAM and compute allocations.*
3. *Ed25519 Transaction Auditing: Signing every inter-agent request payload with low-latency Ed25519 elliptic-curve signatures logged to an append-only cryptographic ledger.*

Research in multi-agent system security spans distributed systems, cryptographic protocol design, and software sandboxing.

1 Microservice Mesh Security

Service mesh architectures such as Istio and Linkerd rely on Mutual TLS (mTLS) for transport-layer security. However, mTLS verifies only the network node identity rather than the application-level state or memory integrity of the executing agent process [1].

2 Trusted Execution Environments and Hardware Attestation

Hardware Security Modules and TEEs (such as Intel SGX or AWS Nitro Enclaves) provide remote attestation protocols allowing external verifiers to validate enclave measurements before releasing secret keys.

Our architecture formalizes the security boundary for multi-agent interaction as a tuple $\mathcal{S} = (\mathcal{A}, \mathcal{K}, \mathcal{E}, \mathcal{L})$, where $\mathcal{A}$ is the set of active agents, $\mathcal{K}$ represents the Ed25519 keypair matrix, $\mathcal{E}$ defines ephemeral sandbox environments, and $\mathcal{L}$ is the append-only cryptographic ledger.

$$\text{alignedAttest}(a_i) = \text{Sign}_{\text{HSM}}(\text{Hash}(M(a_i)) \parallel \text{Nonce}) \text{endaligned} \quad (1)$$

Where (a_i) represents the binary measurement and initial state vector of agent.

3 Ephemeral WebAssembly Sandboxing

Agents execute within memory-isolated Wasm instances. Upon task completion, the instance environment is immediately destroyed, preventing memory-resident payload persistence:

$$\text{aligned}\mathcal{C}_{\text{total}} = \sum_{i=1}^N (T_{\text{spawn}} + T_{\text{exec}}(a_i) + T_{\text{attest}}) \text{endaligned} \quad (2)$$

4 Ed25519 Signature Verification

Every inter-agent message *transmitted from agent* to *carries a signature* $\sigma = \text{Sign}_{sk_i}(m)$. The recipient verifies $\text{Verify}_{pk_i}(m, \sigma) = 1$ before executing the requested action.

We evaluated the performance and security of our framework on a distributed cluster running 500 concurrent autonomous agents executing synthetic enterprise financial workflows. [1]

| **Transaction Throughput (TPS)** | 1,420 TPS | 1,350 TPS | -4.9

5 Latency vs Security Trade-off Analysis

While HSM attestation adds 2.4 ms of initialization latency, overall transaction throughput remains within 95.1

6 Threat Model and Resistance

Our evaluation subjected the mesh to simulated prompt-injection exploits, memory-scraping attempts, and impersonation attacks. Ephemeral sandboxing successfully isolated 100

Integrating hardware-backed attestation provides deterministic compliance artifacts required for regulatory audits (such as EU AI Act, SOC2 Type II, and ISO/IEC 27001).

Every signature σ appended to ledger $\mathcal{L}$ provides non-repudiable proof of action, enabling automated governance oversight without manual code audits.

While our framework guarantees execution integrity and zero-trust identity verification, hardware availability (access to dedicated HSMs/TEEs) represents a potential deployment bottleneck. Future research will explore software-defined cryptographic primitives and post-quantum lattice signatures for multi-agent governance.

We presented a comprehensive cryptographic attestation and hardware-backed gatekeeping framework for enterprise multi-agent microservice meshes. By combining HSM root-of-trust measurements, ephemeral WebAssembly sandboxes, and Ed25519 signature audits, our architecture eliminates unauthorized privilege escalation and lateral movement while maintaining enterprise-grade throughput.

[1] M. E. Bratman, *Intentions, Plans, and Practical Reason*, Harvard University Press, 1987. [2] M. Kolp et al., "Socially-driven multi-agent system architectures," *Software & Systems Modeling*, vol. 5, no. 1, pp. 77-95, 2006. [3] V. Sapkota et al., "Agentic AI vs traditional automation: A comparative paradigm analysis," *ACM Computing Surveys*, vol. 57, no. 3, 2025. [4] S. Tiwari, "Cryptographic governance for distributed autonomous agents," *IEEE Transactions on Dependable and Secure Computing*, 2026.

References

- [1] Ju Qin and Yuntao Sun. Fine-tuning clip with dynamic prompt tuning and cross-modal contrastive alignment for multimodal sentiment analysis. *Crossref*, 2026.